


WEAREDEVELOPERS × BAND · DARK FACTORY

Tablekeeper

built by a three-seat dark factory

One task per stage. No human nudges.
The Reviewer accepts only what it verified
from a fresh clone in a clean, offline container.



2 stages accepted

3 AI seats, 3 model families

4.5 h unattended, \$26.64

3 task messages, 0 nudges

THE CHALLENGE

Build the factory, not the app

- Three or more AI coding seats in one BAND room plan, build and verify a service.
- The human posts one task per stage. No steering, approvals or reruns.
- Every stage must start in a clean container with no network, and keep passing everything earlier stages required.
- Mandates must be generic: no track detail in a seat's standing instructions.

Three seats, three model families

Architect

GLM-5.3

Turns the spec into a numbered ledger; writes work orders, accepts, reports.

Coder

MiniMax-M3

Implements, writes Dockerfile, RUN.m and its own tests; hands off evidence

Reviewer

DeepSeek-V4.1-Flash

Designs probes from the spec in parallel; verifies a fresh clone offline; ACCEPT/REJECT

Engineering the band, not just prompting it

- ACP reply gate: only protocol messages reach the room, so no acknowledgement ping-pong (our first rehearsal burned ~\$156 in a polite loop).
- Turn continuation: a seat that says "starting now" and stops gets a bounded nudge from its own runtime, never from the human.
- Seat shims: human-scoped send commands are refused (one gap found in the run, fixed after it).
- Owner conductor + runaway guard: one task per stage, verified by the official isolated check; seats are pulled out if message rate or spend spikes.

Independent verification, every handoff

- Requirements ledger before code; the ledger, not the shipped tests, defines done.
- Reviewer: fresh clone of the exact revision, clean offline container (2 vCPU / 2 GiB), official harness in isolated mode, its own probes, a diff read.
- 3 rejections in the submitted run, most of them while the shipped checks were already green.
- Escalation after repeated rejects; every verdict is a committed report in reviews/.

SUBMITTED RUN

Results

Stage	Task → REPORT	Minutes	Spend	Owner isolated check
1	05:28 → 06:00	32	\$7.40	claimed 1 · S1 120/120
2	06:01 → 09:01	180	\$15.23	claimed 2 · S1 120/120, S2 25/25
3	09:02 → none	—	—	built, never reviewed: not claimed (owner eval: S1 120/120 · S2 25/25 · S3 7/7)

2 of 4 stages accepted by the band · 4.5 h unattended · \$26.64 model spend · 3 task messages

What we learned the hard way

- Attempt 1 stalled when a seat announced work and stopped → turn continuation.
- Attempt 2 found an earlier run on disk → isolation + "fresh work only" rule.
- Submitted run: the Coder posted one handoff with a human-scoped command, so it shows under the human account → shims now refuse it anywhere.
- Submitted run: a compaction-ended Reviewer turn was never resumed and stage 3 went unreviewed → the gate now continues empty turns too.
- Shipped checks are partial; costs are list price (an upper bound).

TAKE IT WITH YOU

Point the factory at a different problem

- mandates/: three generic seat instructions (scan-clean against the track vocabulary).
- runtime/: reply gate + seat shims, installable as the seats' spawn command.
- factory/: conductor, guard, dispatch builder, recorder, evaluation.
- FACTORY.md: setup, design rationale, measured costs, failure log.
- MIT licensed.